

Deep Learning - Foundations and Concepts

Chapter 4. Single-layer Networks: Regression

nonlineark@github

February 10, 2025

Outline

1 Linear Regression

2 Decision Theory

Basis functions

Consider the linear combinations of fixed nonlinear functions of the input variables:


$$y(x; w) = w_0 + \sum_{m=1}^{M-1} w_m \phi_m(x)$$

where $\phi_m(x)$ are known as basis functions. The parameter w_0 allows for any fixed offset in the data and is sometimes called a bias parameter. If we define $\phi_0(x) = 1$ then $y(x; w)$ becomes:

$$y(x; w) = \sum_{m=0}^{M-1} w_m \phi_m(x) = w^T \phi(x)$$

Basis function

Figure: The linear regression model as a single-layer network



Basis function

Here are some possible choices of basis functions:

- Polynomial: $\phi_m(x) = x^m$.
- Gaussian: $\phi_m(x) = \exp(-\frac{(x-\mu_m)^2}{2s^2})$.
- Sigmoidal: $\phi_m(x) = \frac{1}{1+\exp(-\frac{x-\mu_m}{s})}$.

Maximum likelihood

Consider a data set of inputs $\{x^1, \dots, x^N\}$ with corresponding target values t_1, \dots, t_N . Assume that given the value of x^n , the corresponding value of t_n has a Gaussian distribution. The likelihood function takes the form:

$$p(t_1, \dots, t_N | x^1, \dots, x^N; w, \sigma^2) = \prod_{n=1}^N \mathcal{N}(t_n; w^T \phi(x^n), \sigma^2 I)$$

The negative log of the likelihood function is given by:

$$\begin{aligned} L &= -\log p(t_1, \dots, t_N | x^1, \dots, x^N; w, \sigma^2) \\ &= \frac{N}{2} \log(2\pi) + \frac{N}{2} \log \sigma^2 + \frac{1}{2\sigma^2} \sum_{n=1}^N (t_n - w^T \phi(x^n))^2 \end{aligned}$$

Maximum likelihood

Let's maximize the likelihood function (for simplicity, we will denote $\phi(x^n)$ by ϕ_n):

$$\frac{\partial L}{\partial w} = \frac{1}{\sigma^2} (w^T \sum_{n=1}^N \phi_n \phi_n^T - \sum_{n=1}^N t_n \phi_n^T) = \frac{1}{\sigma^2} (w^T \Phi^T \Phi - t^T \Phi)$$

where:

$$\Phi = (\phi_1 \quad \phi_2 \quad \dots \quad \phi_N)^T$$

We see that:


$$w_{ML} = (\Phi^T \Phi)^{-1} \Phi^T t$$

The quantity $(\Phi^T \Phi)^{-1} \Phi^T$ is known as the Moore-Penrose pseudo-inverse of the matrix Φ . It's easy to calculate σ_{ML}^2 as well:

$$\sigma_{ML}^2 = \frac{1}{N} \sum_{n=1}^N (t_n - w_{ML}^T \phi_n)^2$$

Geometry of least squares

Figure: Geometrical interpretation of the least squares solution



Geometry of least squares

Let Φ_m be the m th column of the matrix Φ , and let $y_{ML} \in \mathbb{R}^N$ be the best approximation to t we obtained by maximizing the likelihood function:

$$y_{ML} = \begin{pmatrix} w_{ML}^T \phi_1 \\ w_{ML}^T \phi_2 \\ \vdots \\ w_{ML}^T \phi_N \end{pmatrix} = \Phi w_{ML} = \sum_{m=0}^{M-1} (w_{ML})_m \Phi_m$$

Here we clearly see that $y_{ML} \in \text{span}(\Phi_0, \dots, \Phi_{M-1})$. In addition, we have:

$$\begin{aligned} \Phi^T y_{ML} &= \Phi^T \Phi w_{ML} = (\Phi^T \Phi)(\Phi^T \Phi)^{-1} \Phi^T t = \Phi^T t \\ (t - y_{ML})^T \Phi &= 0 \quad (t - y_{ML})^T \Phi_m = 0 \end{aligned}$$

That is, $t - y_{ML}$ is orthogonal to each Φ_m , or put another way, y_{ML} is the orthogonal projection of t .

Sequential learning

The maximum likelihood estimator for w involves processing the entire training set in one go. Sometimes we want the data points to be considered one at a time and the model parameters updated after each such presentation. The technique of stochastic (sequential) gradient descent:

- The error function comprises a sum over data points: $E = \sum_n E_n$.
- After presentation of data point n , updates the parameter w using:
$$w^{(\tau+1)} = w^{(\tau)} - \eta \nabla E_n.$$
- τ denotes the iteration number, and η is a training rate parameter.

Sequential learning

For the sum-of-squares error function:

$$E_n = \frac{1}{2}(t_n - w^T \phi_n)^2$$

$$\nabla E_n = -(t_n - w^T \phi_n) \phi_n$$

$$w^{(\tau+1)} = w^{(\tau)} + \eta(t_n - (w^{(\tau)})^T \phi_n) \phi_n$$

Regularized least squares

Adding a regularization term to an error function to control over-fitting:

$$E_D(w) + \lambda E_W(w)$$

For example, if we use the sum-of-squares error function, the total error function becomes:

$$\frac{1}{2} \sum_{n=1}^N (t_n - w^T \phi_n)^2 + \frac{\lambda}{2} w^T w$$

Minimizing this total error function, we obtain:

$$w_{ML} = (\lambda I + \Phi^T \Phi)^{-1} \Phi^T t$$

Multiple outputs

We have considered situations with a single target variable. In some applications, we may wish to predict $K > 1$ target variables. Let's first get the dimensions right:

- There are N input data: x^1, \dots, x^N , where $x^n \in \mathbb{R}^D$.
- There are N target data: t^1, \dots, t^N , where $t^n \in \mathbb{R}^K$.
 - Let $T = (t_1 \ t_2 \ \dots \ t_N)^T \in \mathbb{R}^{N \times K}$
- There is a basis $\phi: \mathbb{R}^D \rightarrow \mathbb{R}^M$, $x \rightarrow \phi(x)$. For simplicity, we denote $\phi(x^n)$ by ϕ_n .
 - Let $\Phi = (\phi_1 \ \phi_2 \ \dots \ \phi_N)^T \in \mathbb{R}^{N \times M}$
- There is a matrix of parameters: $W \in \mathbb{R}^{M \times K}$.

Multiple outputs

Now, let's maximize the likelihood for $y(x; W) = W^T \phi(x)$:

$$\begin{aligned}
 L &= -\log p(t^1, \dots, t^N | x^1, \dots, x^N; W, \sigma^2) \\
 &= -\log \prod_{n=1}^N \mathcal{N}(t^n; W^T \phi_n, \sigma^2 I) \\
 &= \frac{NK}{2} \log(2\pi\sigma^2) + \frac{1}{2\sigma^2} \sum_{n=1}^N \|t^n - W^T \phi_n\|^2 \\
 \frac{\partial L}{\partial W}(W)H &= \frac{1}{\sigma^2} \sum_{n=1}^N (\text{tr}(W^T \phi_n \phi_n^T H) - \text{tr}(t^n \phi_n^T H)) \\
 &= \frac{1}{\sigma^2} \text{tr}((W^T \Phi^T \Phi - T^T \Phi)H) \\
 W_{ML} &= (\Phi^T \Phi)^{-1} \Phi^T T
 \end{aligned}$$

Decision theory

- We have learned from data using maximum likelihood, and the result is a predictive distribution.
- However, for many practical applications we need to predict a specific value.
- In the inference stage, we use the training data to determine a predictive distribution.
- In the decision stage, we use this predictive distribution to determine a specific value.

Decision theory

Problem

Given a predictive distribution $p(t|x)$, determine a specific value $f(x)$, which will be dependent on the input x , that is optimal according to some criterion.

Decision theory

Because we do not know the true value of t , we cannot minimize the loss $L = (f(x) - t)^2$ itself, instead let's minimize the expected loss:

$$E(L) = \iint (f(x) - t)^2 p(x, t) dx dt$$

We want to find $f(x)$ that minimizes $E(L)$:

$$\begin{aligned} \frac{\delta E(L)}{\delta f(x)} &= 2 \int (f(x) - t) p(x, t) dt = 0 \\ f(x) &= \frac{\int t p(x, t) dt}{\int p(x, t) dt} = \frac{\int t p(x, t) dt}{p(x)} = \int t p(t|x) dt = E(t|x) \end{aligned}$$

which is the conditional average of t conditioned on x and is known as the regression function.

Decision theory

Now that we know that the optimal solution is the conditional expectation, we can expand the square term as follows:

$$\begin{aligned}(f(x) - t)^2 &= ((f(x) - E(t|x)) + (E(t|x) - t))^2 \\ &= (f(x) - E(t|x))^2 + 2(f(x) - E(t|x))(E(t|x) - t) + (E(t|x) - t)^2\end{aligned}$$

Decision theory

Let's examine the expectation for each term:

$$\begin{aligned}
 \iint (f(x) - E(t|x))^2 p(x, t) dx dt &= \int (f(x) - E(t|x))^2 \left(\int p(x, t) dt \right) dx \\
 &= \int (f(x) - E(t|x))^2 p(x) dx \\
 \iint (f(x) - E(t|x))(E(t|x) - t) p(x, t) dx dt \\
 &= \int (f(x) - E(t|x)) p(x) \left(\int (E(t|x) - t) p(t|x) dt \right) dx = 0 \\
 \iint (E(t|x) - t)^2 p(x, t) dx dt &= \int p(x) \left(\int (t - E(t|x))^2 p(t|x) dt \right) dx \\
 &= \int \text{var}(t|x) p(x) dx
 \end{aligned}$$

Decision theory

Let's interpret what we have derived here:

$$E(L) = \int (f(x) - E(t|x))^2 p(x) dx + \int \text{var}(t|x) p(x) dx$$

- The first term shows that the optimal least-squares predictor is given by the conditional expectation.
- The second term is the variance of t averaged over x , and represents the intrinsic variability of the target data.