

# Deep Learning - Foundations and Concepts

## Chapter 4. Single-layer Networks: Regression

nonlineark@github

February 9, 2025

# Outline

## 1 Linear Regression

# Basis functions

Consider the linear combinations of fixed nonlinear functions of the input variables:


$$y(x; w) = w_0 + \sum_{m=1}^{M-1} w_m \phi_m(x)$$

where  $\phi_m(x)$  are known as basis functions. The parameter  $w_0$  allows for any fixed offset in the data and is sometimes called a bias parameter. If we define  $\phi_0(x) = 1$  then  $y(x; w)$  becomes:

$$y(x; w) = \sum_{m=0}^{M-1} w_m \phi_m(x) = w^T \phi(x)$$

# Basis function

Figure: The linear regression model as a single-layer network



# Basis function

Here are some possible choices of basis functions:

- Polynomial:  $\phi_m(x) = x^m$ .
- Gaussian:  $\phi_m(x) = \exp(-\frac{(x-\mu_m)^2}{2s^2})$ .
- Sigmoidal:  $\phi_m(x) = \frac{1}{1+\exp(-\frac{x-\mu_m}{s})}$ .

# Maximum likelihood

Consider a data set of inputs  $\{x^1, \dots, x^N\}$  with corresponding target values  $t_1, \dots, t_N$ . Assume that given the value of  $x^n$ , the corresponding value of  $t_n$  has a Gaussian distribution. The likelihood function takes the form:

$$p(t_1, \dots, t_N | x^1, \dots, x^N; w, \sigma^2) = \prod_{n=1}^N \mathcal{N}(t_n; w^T \phi(x^n), \sigma^2 I)$$

The negative log of the likelihood function is given by:

$$\begin{aligned} L &= -\log p(t_1, \dots, t_N | x^1, \dots, x^N; w, \sigma^2) \\ &= \frac{N}{2} \log(2\pi) + \frac{N}{2} \log \sigma^2 + \frac{1}{2\sigma^2} \sum_{n=1}^N (t_n - w^T \phi(x^n))^2 \end{aligned}$$

# Maximum likelihood

Let's maximize the likelihood function (for simplicity, we will denote  $\phi(x^n)$  by  $\phi_n$ ):

$$\frac{\partial L}{\partial w} = \frac{1}{\sigma^2} \left( w^T \sum_{n=1}^N \phi_n \phi_n^T - \sum_{n=1}^N t_n \phi_n^T \right) = \frac{1}{\sigma^2} (w^T \Phi^T \Phi - t^T \Phi)$$

where:

$$\Phi = \begin{pmatrix} \phi_1^T \\ \phi_2^T \\ \vdots \\ \phi_N^T \end{pmatrix}$$

# Maximum likelihood

We see that:

$$w_{ML} = (\Phi^T \Phi)^{-1} \Phi^T t$$


The quantity  $(\Phi^T \Phi)^{-1} \Phi^T$  is known as the Moore-Penrose pseudo-inverse of the matrix  $\Phi$ . It's easy to calculate  $\sigma_{ML}^2$  as well:

$$\sigma_{ML}^2 = \frac{1}{N} \sum_{n=1}^N (t_n - w_{ML}^T \phi_n)^2$$



# Geometry of least squares

Figure: Geometrical interpretation of the least squares solution



# Geometry of least squares

Let  $\Phi_m$  be the  $m$ th column of the matrix  $\Phi$ , and let  $y_{ML} \in \mathbb{R}^N$  be the best approximation to  $t$  we obtained by maximizing the likelihood function:

$$y_{ML} = \begin{pmatrix} w_{ML}^T \phi_1 \\ w_{ML}^T \phi_2 \\ \vdots \\ w_{ML}^T \phi_N \end{pmatrix} = \Phi w_{ML} = \sum_{m=0}^{M-1} (w_{ML})_m \Phi_m$$

Here we clearly see that  $y_{ML} \in \text{span}(\Phi_0, \dots, \Phi_{M-1})$ . In addition, we have:

$$\begin{aligned} \Phi^T y_{ML} &= \Phi^T \Phi w_{ML} = (\Phi^T \Phi)(\Phi^T \Phi)^{-1} \Phi^T t = \Phi^T t \\ (t - y_{ML})^T \Phi &= 0 \quad (t - y_{ML})^T \Phi_m = 0 \end{aligned}$$

That is,  $t - y_{ML}$  is orthogonal to each  $\Phi_m$ , or put another way,  $y_{ML}$  is the orthogonal projection of  $t$ .

# Sequential learning

The maximum likelihood estimator for  $w$  involves processing the entire training set in one go. Sometimes we want the data points to be considered one at a time and the model parameters updated after each such presentation. The technique of stochastic (sequential) gradient descent:

- The error function comprises a sum over data points:  $E = \sum_n E_n$ .
- After presentation of data point  $n$ , updates the parameter  $w$  using:  
$$w^{(\tau+1)} = w^{(\tau)} - \eta \nabla E_n.$$
- $\tau$  denotes the iteration number, and  $\eta$  is a training rate parameter.

# Sequential learning

For the sum-of-squares error function:

$$E_n = \frac{1}{2}(t_n - w^T \phi_n)^2$$

$$\nabla E_n = -(t_n - w^T \phi_n) \phi_n$$

$$w^{(\tau+1)} = w^{(\tau)} + \eta(t_n - (w^{(\tau)})^T \phi_n) \phi_n$$

# Regularized least squares

Adding a regularization term to an error function to control over-fitting:

$$E_D(w) + \lambda E_W(w)$$

For example, if we use the sum-of-squares error function, the total error function becomes:

$$\frac{1}{2} \sum_{n=1}^N (t_n - w^T \phi_n)^2 + \frac{\lambda}{2} w^T w$$

Minimizing this total error function, we obtain:

$$w_{ML} = (\lambda I + \Phi^T \Phi)^{-1} \Phi^T t$$

# Multiple outputs

We have considered situations with a single target variable. In some applications, we may wish to predict  $K > 1$  target variables. Let's first get the dimensions right:

- There are  $N$  input data:  $x^1, \dots, x^N$ , where  $x^n \in \mathbb{R}^D$ .
- There are  $N$  target data:  $t^1, \dots, t^N$ , where  $t^n \in \mathbb{R}^K$ .
  - Let  $T = (t_1 \ t_2 \ \dots \ t_N)^T \in \mathbb{R}^{N \times K}$
- There is a basis  $\phi: \mathbb{R}^D \rightarrow \mathbb{R}^M$ ,  $x \rightarrow \phi(x)$ . For simplicity, we denote  $\phi(x^n)$  by  $\phi_n$ .
  - Let  $\Phi = (\phi_1 \ \phi_2 \ \dots \ \phi_N)^T \in \mathbb{R}^{N \times M}$
- There is a matrix of parameters:  $W \in \mathbb{R}^{M \times K}$ .

# Multiple outputs

Now, let's maximize the likelihood for  $y(x; W) = W^T \phi(x)$ :

$$\begin{aligned}
 L &= -\log p(t^1, \dots, t^N | x^1, \dots, x^N; W, \sigma^2) \\
 &= -\log \prod_{n=1}^N \mathcal{N}(t^n; W^T \phi_n, \sigma^2 I) \\
 &= \frac{NK}{2} \log(2\pi\sigma^2) + \frac{1}{2\sigma^2} \sum_{n=1}^N \|t^n - W^T \phi_n\|^2 \\
 \frac{\partial L}{\partial W}(W)H &= \frac{1}{\sigma^2} \sum_{n=1}^N (\text{tr}(W^T \phi_n \phi_n^T H) - \text{tr}(t^n \phi_n^T H)) \\
 &= \frac{1}{\sigma^2} \text{tr}((W^T \Phi^T \Phi - T^T \Phi)H) \\
 W_{ML} &= (\Phi^T \Phi)^{-1} \Phi^T T
 \end{aligned}$$