

Deep Learning - Foundations and Concepts

Chapter 5. Single-layer Networks: Classification

nonlineark@github

February 14, 2025

Outline

1 Discriminant Functions

2 Decision Theory

3 Generative Classifiers

Discriminant functions

- The goal in classification is to take an input vector $x \in \mathbb{R}^D$ and assign it to one of K discrete classes \mathcal{C}_k .
- A discriminant is a function that takes an input vector x and assigns it to one of K classes, denoted \mathcal{C}_k .
- We will restrict attention to linear discriminants, for which the decision surfaces are hyperplanes.

Two classes


Taking a linear function of the input vector:

$$y(x) = w^T x + w_0$$

- An input vector is assigned to class \mathcal{C}_1 if $y(x) \geq 0$ and to class \mathcal{C}_2 otherwise.
- The decision boundary is a $(D - 1)$ -dimensional hyperplane.

Two classes

Figure: The geometry of a linear discriminant function in two dimensions



Two classes

It's easy to see that:

- w is orthogonal to the decision surface.
- w points to the direction of the increase of y .

Also the value of $y(x)$ gives a signed measure of the perpendicular distance r of the point x from the decision surface:

$$x = x_{\perp} + r \frac{w}{\|w\|}$$

$$y(x) = w^T x + w_0 = w^T x_{\perp} + w_0 + r\|w\| = r\|w\|$$

$$r = \frac{y(x)}{\|w\|}$$

In particular, the signed distance of the origin from the decision surface is given by $\frac{w_0}{\|w\|}$.

Multiple classes

Building a K -class discriminant by combining a number of two-class discriminant functions usually doesn't work:

- One-versus-the-rest.
- One-versus-one.

Multiple classes

Consider a single K -class discriminant comprising K linear functions of the form:

$$y_k(x) = w_k^T x + w_{k0}$$


Assign a point x to class \mathcal{C}_k if $y_k(x) > y_j(x)$ for all $j \neq k$. The decision boundary between class \mathcal{C}_k and \mathcal{C}_j is given by $y_k(x) = y_j(x)$ and corresponds to a $(D - 1)$ -dimensional hyperplane:

$$(w_k - w_j)^T x + (w_{k0} - w_{j0}) = 0$$

The decision regions of such a discriminant are always singly connected and convex.

Multiple classes

Figure: The decision regions for a multi-class linear discriminant



Linear squares for classification

Consider a general classification problem with K classes:

- There are N input data: x^1, \dots, x^N , where $x^n \in \mathbb{R}^D$.
- There are N target data: t^1, \dots, t^N using a 1-of- K binary coding scheme, thus $t^n \in \mathbb{R}^K$.
 - Let $T = (t^1 \quad t^2 \quad \dots \quad t^N)^T \in \mathbb{R}^{N \times K}$.
- Each class \mathcal{C}_k is described by its own linear model so that $y_k(x) = w_k^T x + w_{k0}$.
 - Let $\tilde{w}_k = \begin{pmatrix} w_{k0} \\ w_k \end{pmatrix}$ and $\tilde{W} = (\tilde{w}_1 \quad \tilde{w}_2 \quad \dots \quad \tilde{w}_K) \in \mathbb{R}^{(D+1) \times K}$.
 - Let $\tilde{x} = \begin{pmatrix} 1 \\ x \end{pmatrix}$ and $\tilde{X} = (\tilde{x}^1 \quad \tilde{x}^2 \quad \dots \quad \tilde{x}^N)^T \in \mathbb{R}^{N \times (D+1)}$.
 - Then $y_k(x) = \tilde{w}_k^T \tilde{x}$ and $y(x) = \tilde{W}^T \tilde{x}$.

Linear squares for classification

Let's determine the parameter matrix \tilde{W} by minimizing a sum-of-squares error function:

$$\begin{aligned}
 E_D(\tilde{W}) &= \frac{1}{2} \sum_{i,j} (\tilde{X}\tilde{W} - T)_{ij}^2 \\
 &= \frac{1}{2} \text{tr}((\tilde{X}\tilde{W} - T)^T (\tilde{X}\tilde{W} - T)) \\
 DE_D(\tilde{W})H &= \text{tr}((\tilde{X}\tilde{W} - T)^T \tilde{X}H) \\
 \tilde{W}_* &= (\tilde{X}^T \tilde{X})^{-1} \tilde{X}^T T
 \end{aligned}$$

Linear squares for classification

What property does $y(x) = \tilde{W}_*^T \tilde{x}$ has? Because t^n is using a 1-of- K binary coding scheme, we know:

$$(1 \quad 1 \quad \dots \quad 1) t^n = 1$$

Thus we have:

$$\begin{aligned} (1 \quad 1 \quad \dots \quad 1) y(x) &= (1 \quad 1 \quad \dots \quad 1) \tilde{W}_*^T \tilde{x} \\ &= (1 \quad 1 \quad \dots \quad 1) T^T \tilde{X} (\tilde{X}^T \tilde{X})^{-1} \tilde{x} \\ &= (1 \quad 1 \quad \dots \quad 1) \tilde{X} (\tilde{X}^T \tilde{X})^{-1} \tilde{x} \\ &= \mathbf{e}_1^T \tilde{X}^T \tilde{X} (\tilde{X}^T \tilde{X})^{-1} \tilde{x} = \mathbf{e}_1^T \tilde{x} = 1 \end{aligned}$$

That is, the predictions made by the model will have the property that the elements of $y(x)$ will sum to 1 for any value of x .

Linear squares for classification

- The model outputs cannot be interpreted as probabilities because they are not constrained to lie within the interval $(0, 1)$.
- If the true distribution of the data is markedly different from being Gaussian, the least squares can give poor results.
- Least squares is very sensitive to the presence of outliers (a.k.a., lack robustness).

Misclassification rate

To minimize the chance of assigning x to the wrong class, intuitively we would choose the class having the higher posterior probability.

- Divide the input space into regions \mathcal{R}_k called decision regions.
- All points in \mathcal{R}_k are assigned to class \mathcal{C}_k .

We want to maximize the probability of being correct:

$$p(\text{correct}) = \sum_{k=1}^K p(x \in \mathcal{R}_k, \mathcal{C}_k) = \sum_{k=1}^K \int_{\mathcal{R}_k} p(\mathcal{C}_k|x)p(x)dx$$

It's easy to see that this is maximized when the regions \mathcal{R}_k are chosen such that each x is assigned to the class for which $p(\mathcal{C}_k|x)$ is largest. So the intuition is indeed correct.

Expected loss

- Sometimes, our objective will be more complex than minimizing the number of misclassifications.
- We can introduce a loss function which measure loss incurred in taking any of the available decisions or actions and minimize the total loss.

If the true class for x is \mathcal{C}_k and we assign x to \mathcal{C}_j , we incur some level of loss denoted by L_{kj} . Because we do not know the true class, instead of minimizing the loss function, we minimize its average:

$$E(L) = \sum_k \sum_j \int_{\mathcal{R}_j} L_{kj} p(x, \mathcal{C}_k) dx = \sum_j \int_{\mathcal{R}_j} \sum_k L_{kj} p(\mathcal{C}_k | x) p(x) dx$$

The decision rule that minimizes the expected loss assigns x to the class j for which $\sum_k L_{kj} p(\mathcal{C}_k | x)$ is a minimum.

The reject option

- Classification errors arise when the largest of the posterior probabilities is significantly less than 1.
- Reject option: Avoid making decisions on such cases to obtain a lower error rate.
- Introduce a threshold θ and reject inputs x when the largest of the posterior probabilities is less than or equal to θ :
 - $\theta = 1$: All examples are rejected.
 - $\theta < \frac{1}{K}$: No examples are rejected.

Inference and decision

There are three distinct approaches to solving decision problems:

- Generative models:
 - Solve the inference problem of determining the class-conditional densities $p(x|\mathcal{C}_k)$.
 - Infer the prior class probabilities $p(\mathcal{C}_k)$.
 - Find the posterior class probabilities $p(\mathcal{C}_k|x) = \frac{p(x|\mathcal{C}_k)p(\mathcal{C}_k)}{p(x)}$.
 - Use decision theory to determine the class membership for each new input x .
- Discriminative models:
 - Solve the inference problem of determining the posterior class probabilities $p(\mathcal{C}_k|x)$.
 - Use decision theory to assign each new x to one of the classes.
- Discriminant functions:
 - Find a function that maps each input x directly onto a class label.

Inference and decision

There are many reasons for wanting to compute the posterior probabilities:

- Minimizing risk: What if the loss matrix are subjected to revision from time to time?
- Reject option.
- Compensating for class priors: What if one class occupies 99.9% of the cases (we want a balanced data set to find a more accurate model)?
- Combining models:
 - Combine the outputs of smaller models use the rules of probability.
 - Models can easily be made differentiable with respect to adjustable parameters, which allows them to be composed and trained jointly.

Classifier accuracy

Consider a cancer screening example:

- True positive: The classifier predicts that a person has cancer and is correct.
- False positive (type 1 errors): The classifier predicts that a person has cancer and is wrong.
- True negative: The classifier predicts that a person does not have cancer and is correct.
- False negative (type 2 errors): The classifier predicts that a person does not have cancer and is wrong.

Classifier accuracy

$$\text{Accuracy} = \frac{N_{TP} + N_{TN}}{N_{TP} + N_{FP} + N_{TN} + N_{FN}}$$

$$\text{Precision} = \frac{N_{TP}}{N_{TP} + N_{FP}}$$

$$\text{Recall} = \frac{N_{TP}}{N_{TP} + N_{FN}}$$

$$\text{False positive rate} = \frac{N_{FP}}{N_{FP} + N_{TN}}$$

$$\text{False discovery rate} = \frac{N_{FP}}{N_{FP} + N_{TP}}$$


ROC curve

There is a trade-off between type 1 errors and type 2 errors. To better understand this trade-off, it is useful to plot the ROC (receiver operating characteristic) curve:

- x -axis: False positive rate = $\frac{N_{FP}}{N_{FP} + N_{TN}}$.
- y -axis: True positive rate = $\frac{N_{TP}}{N_{TP} + N_{FN}}$.


ROC curve

Figure: As the decision boundary is moved from ∞ to $-\infty$, the ROC curve is traced out



ROC curve

Figure: The ROC (receiver operating characteristic) curve



ROC curve

Some observations:

- The bottom left corner represents a classifier that always outputs negative.
- The top left corner represents the best possible classifier.
- The top right corner represents a classifier that always outputs positive.
- The diagonal line represents a simple random classifier.

Sometimes it is useful to have a single number that characterises the whole ROC curve:

- The AUC (area under the curve):

- 0.5: Random guessing.
- 1.0: Perfect classifier.

- The F-score: $F = \frac{2 \times \text{precision} \times \text{recall}}{\text{precision} + \text{recall}} = \frac{2N_{TP}}{2N_{TP} + N_{FP} + N_{FN}}.$

Activation functions

In linear regression, the model prediction is given by:

$$y(x; w) = w^T x + w_0$$

which gives a continuous-valued output in the range $(-\infty, +\infty)$. For classification problems, we wish to predict posterior probabilities in the range $(0, 1)$, which could be achieved using an activation function:

$$y(x; w) = f(w^T x + w_0)$$

We see that the decision surfaces are linear functions of x . For this reason, these models are called generalized linear models.

Activation functions

For two classes, we can use the logistic sigmoid function $\sigma(a) = \frac{1}{1+\exp(-a)}$ as the activation function:

$$p(\mathcal{C}_1|x) = \sigma(a(x)) = \frac{1}{1 + \exp(-a(x))}$$

Compare with:

$$\begin{aligned} p(\mathcal{C}_1|x) &= \frac{p(x|\mathcal{C}_1)p(\mathcal{C}_1)}{p(x|\mathcal{C}_1)p(\mathcal{C}_1) + p(x|\mathcal{C}_2)p(\mathcal{C}_2)} \\ &= \frac{1}{1 + \frac{p(x|\mathcal{C}_2)p(\mathcal{C}_2)}{p(x|\mathcal{C}_1)p(\mathcal{C}_1)}} \end{aligned}$$

We see that:

$$a(x) = \log \frac{p(x|\mathcal{C}_1)p(\mathcal{C}_1)}{p(x|\mathcal{C}_2)p(\mathcal{C}_2)}$$

Activation functions

The softmax function is defined by:

$$\text{softmax}(a) = \frac{1}{\sum_{k=1}^K \exp(a_k)} \begin{pmatrix} \exp(a_1) \\ \vdots \\ \exp(a_K) \end{pmatrix}$$

For multiple classes, we can use the softmax function as the activation function:

$$\begin{pmatrix} p(\mathcal{C}_1|x) \\ \vdots \\ p(\mathcal{C}_K|x) \end{pmatrix} = \text{softmax}(a_1(x), \dots, a_K(x))$$

Compare with:

$$p(\mathcal{C}_k|x) = \frac{p(x|\mathcal{C}_k)p(\mathcal{C}_k)}{\sum_j p(x|\mathcal{C}_j)p(\mathcal{C}_j)}$$

We see that:

$$a_k(x) = \log(p(x|\mathcal{C}_k)p(\mathcal{C}_k))$$

Continuous inputs

Let's assume that the class-conditional densities are Gaussian. To start with, we will assume that all classes share the same covariance matrix Σ :

$$p(x|\mathcal{C}_k) = \frac{1}{(2\pi)^{\frac{D}{2}} (\det \Sigma)^{\frac{1}{2}}} \exp\left(-\frac{1}{2}(x - \mu_k)^T \Sigma^{-1}(x - \mu_k)\right)$$

We will see that this lead to generalized linear models.

Continuous inputs

For two classes:

$$p(\mathcal{C}_1|x) = \sigma(a(x))$$

where:

$$a(x) = w^T x + w_0$$

$$w = \Sigma^{-1}(\mu_1 - \mu_2)$$

$$w_0 = -\frac{1}{2}\mu_1^T \Sigma^{-1} \mu_1 + \frac{1}{2}\mu_2^T \Sigma^{-1} \mu_2 + \log \frac{p(\mathcal{C}_1)}{p(\mathcal{C}_2)}$$

Continuous inputs

For multiple classes:

$$\begin{pmatrix} p(\mathcal{C}_1|x) \\ \vdots \\ p(\mathcal{C}_K|x) \end{pmatrix} = \text{softmax}(a_1(x), \dots, a_K(x))$$

where:

$$a_k(x) = w_k^T x + w_{k0}$$

$$w_k = \Sigma^{-1} \mu_k$$

$$w_{k0} = -\frac{1}{2} \mu_k^T \Sigma^{-1} \mu_k + \log p(\mathcal{C}_k)$$

Maximum likelihood solution

- There are N input data: x^1, \dots, x^N , where $x^n \in \mathbb{R}^D$.
- There are N target data: t^1, \dots, t^N using a 1-of- K binary coding scheme, thus $t^n \in \mathbb{R}^K$.
- The total number of data points in class \mathcal{C}_k is denoted by N_k .
- Prior class probabilities are denoted by $\pi_k = p(\mathcal{C}_k)$.
- Class-conditional densities are Gaussian: $p(x|\mathcal{C}_k) = \mathcal{N}(x; \mu_k, \Sigma)$.

Maximum likelihood solution

$$\begin{aligned}
 p(t^n|x^n) &= \prod_{k=1}^K p(\mathcal{C}_k|x^n)^{t_k^n} = \prod_{k=1}^K \left(\frac{p(x^n|\mathcal{C}_k)\pi_k}{p(x^n)} \right)^{t_k^n} \\
 &= \frac{1}{p(x^n)} \prod_{k=1}^K \pi_k^{t_k^n} \prod_{k=1}^K \mathcal{N}(x^n; \mu_k, \Sigma)^{t_k^n} \\
 -\log p(t^n|x^n) &= -\sum_{k=1}^K t_k^n \log \pi_k \\
 &\quad + \frac{1}{2} \log \det \Sigma + \frac{1}{2} \sum_{k=1}^K t_k^n (x^n - \mu_k)^T \Sigma^{-1} (x^n - \mu_k) \\
 &\quad + \log p(x^n) + \frac{D}{2} \log 2\pi
 \end{aligned}$$

Maximum likelihood solution

$$\begin{aligned}
 L &= -\log p(t^1, \dots, t^N | x^1, \dots, x^N) = -\sum_{n=1}^N \log p(t^n | x^n) \\
 &= -\sum_{k=1}^K N_k \log \pi_k \\
 &\quad + \frac{N}{2} \log \det \Sigma + \frac{1}{2} \sum_{k=1}^K \sum_{n=1}^N t_k^n (x^n - \mu_k)^T \Sigma^{-1} (x^n - \mu_k) \\
 &\quad + \sum_{n=1}^N \log p(x^n) + \frac{ND}{2} \log 2\pi
 \end{aligned}$$

Maximum likelihood solution

$$\frac{\partial L}{\partial \pi_k} = \frac{N_K}{\pi_K} - \frac{N_k}{\pi_k} \quad \pi_k = \frac{N_k}{N}$$

$$\frac{\partial L}{\partial \mu_k} = (N_k \mu_k - \sum_{x^n \in \mathcal{C}_k} x^n)^T \Sigma^{-1} \quad \mu_k = \frac{1}{N_k} \sum_{x^n \in \mathcal{C}_k} x^n$$

$$\frac{\partial L}{\partial \Lambda}(\Lambda)H = \frac{1}{2} \text{tr}((\sum_{k=1}^K \sum_{x^n \in \mathcal{C}_k} (x^n - \mu_k)(x^n - \mu_k)^T - N\Sigma)H)$$

$$\Sigma = \sum_{k=1}^K \frac{N_k}{N} S_k \quad S_k = \frac{1}{N_k} \sum_{x^n \in \mathcal{C}_k} (x^n - \mu_k)(x^n - \mu_k)^T$$

Discrete features

Suppose $x \in \mathbb{R}^D$ is a feature vector, where each feature $x_d \in \{0, 1\}$. And further suppose that the different features are independent when conditioned on the class \mathcal{C}_k . So we have:

$$p(x|\mathcal{C}_k) = \prod_{d=1}^D p(x_d|\mathcal{C}_k) = \prod_{d=1}^D \mu_{dk}^{x_d} (1 - \mu_{dk})^{1-x_d}$$

Using a softmax activation function, we see that:

$$a_k(x) = \sum_{d=1}^D (x_d \log \mu_{dk} + (1 - x_d) \log(1 - \mu_{dk})) + \log p(\mathcal{C}_k)$$

which again are linear functions of the input values x .

Exponential family

If the class-conditional densities $p(x|\mathcal{C}_k)$ are members of the subset of the exponential family of distributions given by:

$$p(x|\mathcal{C}_k; \lambda_k, s) = \frac{1}{s} h\left(\frac{1}{s}x\right) g(\lambda_k) \exp\left(\frac{1}{s} \lambda_k^T x\right)$$

the resulting model will be a generalized linear model. For two classes using the logistic sigmoid activation function:

$$a(x) = \frac{1}{s} (\lambda_1 - \lambda_2)^T x + \log \frac{g(\lambda_1)}{g(\lambda_2)} + \log \frac{p(\mathcal{C}_1)}{p(\mathcal{C}_2)}$$

For multiple classes using the softmax activation function:

$$a_k(x) = \frac{1}{s} \lambda_k^T x + \log g(\lambda_k) + \log p(\mathcal{C}_k)$$