

# Deep Learning - Foundations and Concepts

## Chapter 19. Autoencoders

nonlineark@github

April 19, 2025


# Outline

1 Deterministic Autoencoders

2 Variational Autoencoders

# Linear autoencoders

Figure: An autoencoder neural network having two layers of weights



# Linear autoencoders

Consider a two-layer neural network having  $D$  inputs,  $D$  output units and  $M$  hidden units, with  $M < D$ . The targets used to train the network are simply the input vectors themselves, so that the network attempts to map each input vector onto itself. We choose a sum-of-squares error of the form:


$$E(w) = \frac{1}{2} \sum_{n=1}^N \|y(x_n; w) - x_n\|^2$$

# Linear autoencoders

- If the hidden units have linear activation functions, then it can be shown that when the error function is minimized, the network performs a projection onto the  $M$ -dimensional subspace that is spanned by the first  $M$  principal components of the data.
- Even with nonlinear hidden units, the minimum error solution is again given by the projection onto the principal component subspace. There is therefore no advantage in using two-layer neural networks to perform dimensionality reduction.

# Deep autoencoders

**Figure:** A four-layer auto-associative network that can perform a nonlinear dimensionality reduction



# Deep autoencoders

Consider the four-layer auto-associative network shown on the previous slide:

- The output units are linear, and the  $M$  units in the second layer can also be linear.
- However, the first and third layers have sigmoidal nonlinear activation functions.

We can view this network as two successive functional mappings  $F_1$  and  $F_2$ :

- $F_1$  projects the original  $D$ -dimensional data onto an  $M$ -dimensional subspace defined by the activations of the units in the second layer.
- $F_2$  maps from the  $M$ -dimensional hidden space back into the original  $D$ -dimensional input space.

# Deep autoencoders

- Such a network effectively performs a nonlinear form of PCA.
- However:
  - Training the network now involves a nonlinear optimization, and computationally intensive nonlinear optimization techniques must be used.
  - There is the risk of finding a sub-optimal local minimum of the error function.
  - The dimensionality of the subspace must be specified before training the network.



# Sparse autoencoders

Instead of limiting the number of nodes in one of the hidden layers in the network, an alternative way to constrain the internal representation is to use a regularizer to encourage a sparse representation:

$$\tilde{E}(w) = E(w) + \lambda \sum_{k=1}^K |z_k|$$

where  $E(w)$  is the unregularized error, and the sum over  $k$  is taken over the activation values of all the units in one of the hidden layers.

# Denoising autoencoders

The idea of denoising autoencoders is to take each input vector  $x_n$  and to corrupt it with noise to give a modified vector  $\tilde{x}_n$  which is then input to an autoencoder to give an output  $y(\tilde{x}_n; w)$ . The network is trained to reconstruct the original noise-free input vector:

$$E(w) = \sum_{n=1}^N ||y(\tilde{x}_n; w) - x_n||^2$$

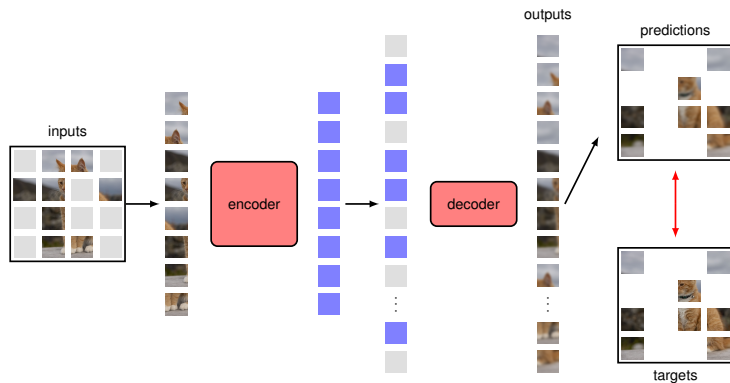
# Denoising autoencoders

- One form of noise involves setting a randomly chosen subset of the input variables to zero.
- An alternative approach is to add independent zero-mean Gaussian noise to every input variable, where the scale of the noise is set by the variance of the Gaussian.

# Masked autoencoders

- In a masked autoencoder, a deep network is used to reconstruct an image given a corrupted version of that image as input. The form of corruption is masking, or dropping out, part of the input image.
- This technique is generally used in combination with a vision transformer architecture.
- Compared to language, images have much more redundancy along with strong local correlations. The best internal representations are learned when a relatively high proportion of the input image is masked, typically 75%.
- The decoder is discarded and the encoder is applied to the full image with no masking and with a fresh set of output layers that are fine-tuned for the required application.

# Masked autoencoders



# Variational autoencoders

Because evaluating the likelihood function for a latent-variable model is intractable, the variational autoencoders (VAEs) instead work with an approximation to this likelihood when training the model:

- Use of the evidence lower bound (ELBO) to approximate the likelihood function.
- Amortized inference in which a second model, the encoder network, is used to approximate the posterior distributions over latent variables in the E step.
- Making the training of the encoder model tractable using the reparameterization trick.

# Evidence lower bound

Consider a generative model:

- The distribution over the  $M$ -dimensional latent variable  $z$  is given by a zero-mean unit-variance Gaussian:  $p(z) = \mathcal{N}(z; 0, I)$ .
- The conditional distribution  $p(x|z; w)$  over the  $D$ -dimensional data variable  $x$  is governed by the output of a deep neural network  $g(z; w)$ .

# Evidence lower bound

We know that for an arbitrary probability distribution  $q(z)$  over a space described by the latent variable  $z$ , we have:

$$\log p(x; w) = \mathcal{L}(w) + \text{KL}(q(z) || p(z|x; w)) \geq \mathcal{L}(w)$$

$$\mathcal{L}(w) = \int q(z) \log \frac{p(x|z; w)p(z)}{q(z)} dz$$

$$\text{KL}(q(z) || p(z|x; w)) = - \int q(z) \log \frac{p(z|x; w)}{q(z)} dz$$

Although the log likelihood  $\log p(x; w)$  is intractable, we will seek a way to evaluate  $\mathcal{L}(w)$  as an approximation to the true log likelihood.



# Evidence lower bound

Now consider a set of training data points  $\mathcal{D} = \{x_1, \dots, x_N\}$ , which are assumed to be drawn independently from the model distribution  $p(x)$ . The log likelihood function for this data set is given by:

$$\log p(\mathcal{D}; w) = \sum_{n=1}^N \mathcal{L}_n(w) + \sum_{n=1}^N \text{KL}(q_n(z_n) || p(z_n | x_n; w))$$
$$\mathcal{L}_n(w) = \int q_n(z_n) \log \frac{p(x_n | z_n; w) p(z_n)}{q_n(z_n)} dz_n$$

# Evidence lower bound

- Theoretically, we can set each  $q_n(z_n)$  equal to the corresponding posterior distribution  $p(z_n|x_n; w)$ , which gives zero Kullback-Leibler divergence, and hence the lower bound is equal to the true log likelihood.
- However, in practice, exact evaluation of  $p(z_n|x_n; w)$  is intractable. We therefore need to find an approximation to the posterior distribution.

# Amortized inference

- In the variational autoencoder, we train a single neural network, called the encoder network, to approximate all the posterior distributions  $p(z_n|x_n; w)$ .
- The encoder produces a single distribution  $q(z|x; \phi)$  that is conditioned on  $x$ , where  $\phi$  represents the parameters of the network.
- The objective function, given by the evidence lower bound, now has a dependence on  $\phi$  as well as  $w$ , and we maximize the bound jointly with respect to both sets of parameters.


# Amortized inference

A typical choice for the encoder is a Gaussian distribution with a diagonal covariance matrix:

$$q(z|x; \phi) = \prod_{m=1}^M \mathcal{N}(z_m; \mu_m(x, \phi), \sigma_m^2(x, \phi))$$


# Amortized inference

Figure: Illustration of the optimization of the ELBO



# Amortized inference

Figure: Comparison of the EM algorithm with ELBO optimization in a VAE



# The reparameterization trick

For data point  $x_n$ , we can write the contribution to the lower bound in the form:

$$\begin{aligned}\mathcal{L}(w, \phi) &= \int q(z_n|x_n; \phi) \log \frac{p(x_n|z_n; w)p(z_n)}{q(z_n|x_n; \phi)} dz_n \\ &= \int q(z_n|x_n; \phi) \log p(x_n|z_n; w) dz_n - \text{KL}(q(z_n|x_n; \phi) || p(z_n))\end{aligned}$$

# The reparameterization trick

For the first term, we could try to approximate the intergral over  $z_n$  with a simple Monte Carlo estimator:

$$\int q(z_n|x_n; \phi) \log p(x_n|z_n; w) dz_n \approx \frac{1}{L} \sum_{l=1}^L \log p(x_n|z_n^{(l)}; w)$$

where  $z_n^{(l)}$  are samples drawn from the encoder distribution  $q(z_n|x_n; \phi)$ . This is easily differentiated with respect to  $w$ , but the gradient with respect to  $\phi$  is problematic.



# The reparameterization trick

We can resolve this by making use of the reparameterization trick in which we reformulate the Monte Carlo sampling procedure such that derivatives with respect to  $\phi$  can be calculated explicitly. Instead of drawing samples of  $z_n$  directly, we draw samples from  $\mathcal{N}(\epsilon; 0, I)$ :

$$z_{nm}^{(l)} = \mu_{nm} + \sigma_{nm}\epsilon_{nm}^{(l)}$$

where  $\mu_{nm} = \mu_m(x_n, \phi)$  and  $\sigma_{nm} = \sigma_m(x_n, \phi)$ . This makes the dependence on  $\phi$  explicit and allows gradients with respect to  $\phi$  to be evaluated.

# The reparameterization trick

The second term for  $\mathcal{L}(w, \phi)$  is a Kullback-Leibler divergence between two Gaussian distributions and can be evaluated analytically:

$$\text{KL}(q(z_n|x_n; \phi) || p(z_n)) = \frac{1}{2} \sum_{m=1}^M (-1 - \log \sigma_{nm}^2 + \mu_{nm}^2 + \sigma_{nm}^2)$$

The full error function for the VAE, therefore becomes:

$$\begin{aligned} z_n^{(l)} &= \mu_n + \text{diag}(\sigma_{n1}, \dots, \sigma_{nM}) \epsilon_n^{(l)} \\ \mathcal{L}(w, \phi) &= \sum_{n=1}^N \left( \frac{1}{L} \sum_{l=1}^L \log p(x_n | z_n^{(l)}; w) \right. \\ &\quad \left. + \frac{1}{2} \sum_{m=1}^M (1 + \log \sigma_{nm}^2 - \mu_{nm}^2 - \sigma_{nm}^2) \right) \end{aligned}$$

# The reparameterization trick

---

## Algorithm 1: Variational autoencoder training

---

**repeat**

$\mathcal{L} \leftarrow 0;$

**for**  $m \leftarrow 1$  **to**  $M$  **do**

$\epsilon_{nm} \sim \mathcal{N}(\epsilon; 0, 1);$

$z_{nm} \leftarrow \mu_{nm} + \sigma_{nm}\epsilon_{nm};$

$\mathcal{L} \leftarrow \mathcal{L} + \frac{1}{2}(1 + \log \sigma_{nm}^2 - \mu_{nm}^2 - \sigma_{nm}^2);$

**end**

$\mathcal{L} \leftarrow \mathcal{L} + \log p(x_n | z_n; w);$

$w \leftarrow w + \eta \nabla_w \mathcal{L};$

$\phi \leftarrow \phi + \eta \nabla_\phi \mathcal{L};$

**until** *converged*;

**return**  $w, \phi;$

---

# The reparameterization trick

Problems when training VAEs:

- Posterior collapse: The variational distribution  $q(z|x; \phi)$  converges to the prior distribution  $p(z)$  and therefore becomes uninformative because it no longer depends on  $x$ .
- Latent code is not compressed.

Both problems can be addressed by introducing a coefficient  $\beta$  in front of the second term in  $\mathcal{L}(w, \phi)$  to control the regularization effectiveness of the Kullback-Leibler divergence (further references).